

- 9-1 Using Group By and Having Clauses

HAVING	Used to specify which groups are to be displayed; restricts groups that do not meet group criteria
GROUP BY	Divides the rows in a table into groups

In the SQL query shown below, which of the following is true about this query?

TRUE a. Kimberly Grant would not appear in the results set.

FALSE b. The GROUP BY clause has an error because the manager_id is not listed in the SELECT clause.

FALSE c. Only salaries greater than 16001 will be in the result set.

FALSE d. Names beginning with Ki will appear after names beginning with Ko.

FALSE e. Last names such as King and Kochhar will be returned even if they don't have salaries > 16000.

```
SELECT last_name, MAX(salary)
FROM employees
WHERE last_name LIKE 'K%'
GROUP BY manager_id, last_name
HAVING MAX(salary) > 16000
ORDER BY last_name DESC ;
```

Each of the following SQL queries has an error. Find the error and correct it. Use Oracle Application Express to verify that your corrections produce the desired results.

a. SELECT manager_id, **AVG(SALARY)**
FROM employees
WHERE AVG(salary) < 16000
GROUP BY manager_id;

b. SELECT cd_number, COUNT(*)
FROM d_cds
WHERE cd_number < 93;

c. SELECT type_code, MAX(TO_NUMBER(REPLACE(duration, ' min', ''))) || ' min' as
"max duration"

FROM d_songs

WHERE duration IN('3 min', '6 min', '10 min')

AND HAVING ID < 50

GROUP by type_code;

d. SELECT loc_type, AVG(

CASE

WHEN INSTR(rental_fee, '/hour') != 0 THEN

TO_NUMBER(REPLACE(rental_fee, '/hour', ''))*5

WHEN INSTR(rental_fee, '/flat fee') != 0 THEN

TO_NUMBER(REPLACE(rental_fee, '/flat fee', ''))

WHEN INSTR(rental_fee, '/per person') != 0 THEN

TO_NUMBER(REPLACE(rental_fee, '/per person', ''))*10

ELSE 0

END

)

AS Fee

FROM d_venues

WHERE id < 100

GROUP BY loc_type

ORDER BY 2;

3. Rewrite the following query to accomplish the same result:

SELECT track, MAX(song_id)

FROM d_track_listings

WHERE track IN (1, 2, 3)

GROUP BY track;

4. Indicate True or False

TRUE a. If you include a group function and any other individual columns in a SELECT clause, then each individual column must also appear in the GROUP BY clause.

FALSE b. You can use a column alias in the GROUP BY clause.

FALSE c. The GROUP BY clause always includes a group function.

5. Write a query that will return both the maximum and minimum average salary grouped by department from the employees table.

```
SELECT ROUND(MAX(AVG(salary)),2) as "Maximum Average of Departments",  
ROUND(MIN(AVG(salary)),2) "Minimum Average of Departments"  
FROM employees  
GROUP BY department_id;
```

6. Write a query that will return the average of the maximum salaries in each department for the employees table.

```
SELECT AVG(MAX(salary))  
  
FROM employees  
  
GROUP BY department_id;
```

- 9-2 Using Rollup and Cube Operations, and Grouping Sets

ROLLUP	Used to create subtotals that roll up from the most detailed level to a grand total, following a grouping list specified in the clause
CUBE	An extension to the GROUP BY clause like ROLLUP that produces cross-tabulation reports
GROUPING SETS	Used to specify multiple grouping of data

Try It / Solve It

1. Within the Employees table, each manager_id is the manager of one or more employees who each have a job_id and earn a salary. For each manager, what is the total salary earned by all of the employees within each job_id? Write a query

to display the Manager_id, job_id, and total salary. Include in the result the subtotal salary for each manager and a grand total of all salaries.

```
SELECT manager_id, job_id, SUM(salary) "total salary",  
GROUPING(manager_id), GROUPING(job_id)  
FROM employees  
GROUP BY ROLLUP(manager_id, job_id);
```

2. Amend the previous query to also include a subtotal salary for each job_id regardless of the manager_id.

```
SELECT manager_id, job_id, SUM(salary) "total salary",  
GROUPING(manager_id), GROUPING(job_id)  
FROM employees  
GROUP BY CUBE(manager_id, job_id);
```

3. Using GROUPING SETS, write a query to show the following groupings:

- department_id, manager_id, job_id
- manager_id, job_id
- department_id, manager_id

```
SELECT department_id, manager_id, job_id, SUM(salary) "total salary",  
GROUPING(department_id), GROUPING(manager_id), GROUPING(job_id)
```

```
FROM employees
```

```
GROUP BY GROUPING SETS((department_id, manager_id, job_id), (manager_id,  
job_id), (department_id, manager_id));
```

- 9-3 Using Set Operators

UNION	Operator that returns all rows from both tables and eliminates duplicates
TO_CHAR(NULL) or TO_DATE(NULL) or TO_NUMBER(NULL)	Columns that were made up to match queries in another table that are not in both tables
UNION ALL	Operator that returns all rows from both tables, including duplicates
Set operators	Used to combine results into one single result from multiple SELECT statements
MINUS	Operator that returns rows that are unique to each table
INTERSECT	Operator that returns rows common to both tables

1. Name the different Set operators?

UNION

UNION ALL

MINUS

INTERSECT

2. Write one query to return the employee_id, job_id, hire_date, and department_id of all employees and a second query listing employee_id, job_id, start_date, and department_id from the job_history table and combine the results as one single output. Make sure you suppress duplicates in the output.

```
SELECT employee_id, job_id, hire_date, department_id
```

```
FROM employees
```

```
UNION
```

```
SELECT employee_id, job_id, start_date, department_id
```

```
FROM job_history
```

```
ORDER BY employee_id, hire_date;
```

3. Amend the previous statement to not suppress duplicates and examine the output. How many extra rows did you get returned and which were they? Sort the output by employee_id to make it easier to spot.

```
SELECT employee_id, job_id, hire_date, department_id
FROM employees
UNION ALL
SELECT employee_id, job_id, start_date, department_id
FROM job_history
ORDER BY employee_id, hire_date;
```

4. List all employees who have not changed jobs even once. (Such employees are not found in the job_history table)

```
SELECT DISTINCT employee_id
FROM employees
MINUS
SELECT DISTINCT employee_id
FROM job_history;
```

5. List the employees that HAVE changed their jobs at least once.

```
SELECT DISTINCT employee_id
FROM employees
INTERSECT
SELECT DISTINCT employee_id
FROM job_history;
```

6. Using the UNION operator, write a query that displays the employee_id, job_id, and salary of ALL present and past employees. If a salary is not found, then just display a 0 (zero) in its place.

```
SELECT employee_id, job_id, NVL(salary, 0)
FROM employees
UNION
SELECT employee_id, job_id, 0
FROM job_history
ORDER BY employee_id;
```

